

Sentence Embedding Models for Ancient Greek Using Multilingual Knowledge Distillation

Kevin Krahn and Derrick Tate and Andrew C. Lamicela

Sattler College, Boston, MA

{kevin.krahn24, dtate, alamicela}@sattler.edu

Abstract

Contextual language models have been trained on Classical languages, including Ancient Greek and Latin, for tasks such as lemmatization, morphological tagging, part of speech tagging, authorship attribution, and detection of scribal errors. However, high-quality sentence embedding models for these historical languages are significantly more difficult to achieve due to the lack of training data. In this work, we use a multilingual knowledge distillation approach to train BERT models to produce sentence embeddings for Ancient Greek text. The state-of-the-art sentence embedding approaches for high-resource languages use massive datasets, but our distillation approach allows our Ancient Greek models to inherit the properties of these models while using a relatively small amount of translated sentence data. We build a parallel sentence dataset using a sentence-embedding alignment method to align Ancient Greek documents with English translations, and use this dataset to train our models. We evaluate our models on translation search, semantic similarity, and semantic retrieval tasks and investigate translation bias. We make our training and evaluation datasets freely available at [this url](#).

1 Introduction

Sentence embedding models, which map sentences or other sequences of text to a dense vector space, such that semantically similar sentences are close together in the vector space, have many applications in NLP. Current state-of-the-art sentence embedding models, however, are trained on modern, high-resource languages such as English and use massive datasets consisting of billions of sentence pairs (Ni et al., 2022). A different approach is needed for historical languages, which have much less data available.

In this work, we train several sentence embedding models for Ancient Greek. Many more Ancient Greek texts have survived compared to texts

from most other ancient languages, which makes sentence embedding models both more feasible and useful.

Several previous works have trained language models for Ancient Greek. Johnson et al. (2021) introduced the Classical Language Toolkit (CLTK) which includes several tools for Ancient Greek processing, including static word embeddings. Singh et al. (2021) fine-tuned a Modern Greek BERT model (Koutsikakis et al., 2020) on Ancient Greek text for PoS tagging, morphological tagging, and lemmatization tasks. Yamshchikov et al. (2022) trained a BERT model for authorship classification of Pseudo-Plutarch texts. Cowen-Breen et al. (2023) trained another BERT model for the purpose of identifying errors in scribal transmission. Riemenschneider and Frank (2023) produced the most comprehensive work on Classical language models to date, training multiple models on a large multilingual corpus of Ancient Greek, Latin, and English texts and comprehensively evaluating and comparing their new models to previous models on a variety of tasks. None of these works, however, produce sentence embedding models for Ancient Greek.

Although there are many digitized Ancient Greek texts available, there is a lack of suitable training data for training sentence embedding models from scratch. The best approaches for high-resource languages involve large human-annotated datasets, such as the natural language inference (NLI) datasets used by Sentence-BERT (Reimers and Gurevych, 2019). Needless to say, such datasets are not available for Ancient Greek.

Following Reimers and Gurevych (2020), we use *multilingual knowledge distillation* to train sentence embedding models with an aligned vector space for Ancient Greek and English. Given a teacher model M for a language s , and a dataset of translated sentences $((s_1, t_1) \dots (s_n, t_n))$ where s_i and t_i are parallel sentences, we train a new stu-




Figure 1: Multilingual knowledge distillation for English to Ancient Greek sentence pairs.

dent model \hat{M} to mimic the sentence embeddings of the teacher M using mean squared loss, such that $\hat{M}(s_i) \approx M(t_i)$ and $\hat{M}(t_i) \approx M(s_i)$. In our case, the teacher model is English and the student model learns both Greek¹ and English embeddings.

This approach has numerous advantages: 1) it requires a relatively small amount of training data, 2) the student model inherits the vector space properties of a state-of-the-art English sentence embedding model, 3) the student model is multilingual, and 4) the vector spaces are aligned across languages.

The cross-lingual nature of this approach is especially useful for Ancient Greek semantic retrieval, since it is much easier to formulate search queries in English than in Ancient Greek. Although it is possible to operate on the English translations of Greek texts, translations are not readily available for all Greek texts, and the available translations are usually not aligned at the sentence level, making it difficult to quickly find the corresponding Greek text. Furthermore, English translations can suffer from various kinds of translator bias, whereas a language model that operates directly on the Greek text can offer an “average” of multiple translators’ interpretations of the text (See Section 4.4).

We produce a training dataset of parallel sentences using a two-step translation alignment process: an initial, smaller dataset was produced using a sentence-length heuristic and dictionary-based alignment technique (Halász et al., 2007), and this initial dataset was used to train an intermediate multilingual sentence embedding model, which was used to align a larger dataset using the approach introduced by Liu and Zhu (2023), which

¹When we refer to “Greek” in an unqualified way in this paper we are referring to Ancient Greek.

uses sentence embeddings for state-of-the-art alignment quality.

We create new evaluation datasets for Ancient Greek translation search, semantic textual similarity (STS), and semantic retrieval (SR) and we evaluate our models on these datasets.

In summary, our contributions are as follows:

1. We use a multilingual knowledge distillation approach to train several Ancient Greek sentence embedding models.
2. We use translation alignment to produce a dataset of Ancient Greek sentences and their English translations.
3. We develop evaluation datasets for translation search, semantic retrieval, and semantic textual similarity, and we evaluate our sentence embedding models on these tasks.

2 Training

2.1 Base Models

To train a sentence embedding model, we first need a base language model trained on Ancient Greek text. The existing Ancient Greek language models were unsuitable for our purposes; most of them are monolingual, but we are training a multilingual model. The models trained by Riemenschneider and Frank (2023) would be the best candidates because they include English, but one of their goals was to avoid contamination from modern languages, such as modern concepts and technology like cellphones which were unknown in antiquity. However, for us this is not a concern, since one of our goals is to train a model to facilitate semantic search with modern language and terminology.

Instead, we fine-tune multilingual BERT-base (mBERT) (Devlin et al., 2019) and XLM-RoBERTa-base (XLM-R) (Conneau et al., 2020) for our base models. Pires et al. (2019) shows that **low-resource languages can benefit from multilingual pre-training**. We use masked language modeling (MLM) to fine-tune mBERT, (denoted as GRCmBERT) and XLM-R (denoted as GRCXLM-R) with Ancient Greek text, and we use these as base models. See Appendix A for training details.

Both mBERT and XLM-R were trained on Modern Greek, among many other languages, but not on Ancient Greek, and hence one **disadvantage of these models is that their tokenizers are not optimized for Ancient Greek morphology**, which could

Model	Symbols/token	Words/token
mBERT	2.29	0.37
XLM-R	2.66	0.43

Table 1: The XLM-R tokenizer produces longer tokens and a higher number of words per token on Ancient Greek text compared to the mBERT tokenizer.

negatively impact performance (Park et al., 2021; Hofmann et al., 2021).

We use a similar approach to Yamshchikov et al. (2022) to compare the mBERT and XLM-R tokenizers. We take a random sample of 20k Ancient Greek sentences from the pre-training corpus and compute the average token length and average words per token for a rough estimation of tokenization quality (See Table 1). The XLM-R tokenizer scores higher on both metrics compared to the mBERT tokenizer. However, a higher score for either metric does not guarantee superior performance in downstream tasks, since it does not measure how well the sub-word tokens capture Ancient Greek morphology.

2.2 Knowledge Distillation

To train multilingual sentence embedding models on English and Ancient Greek with an aligned vector space we use *multilingual knowledge distillation* (Reimers and Gurevych, 2020). This process requires a teacher model M for a source language s , and a dataset of translated sentences $((s_1, t_1), \dots, (s_n, t_n))$ where s_i and t_i are parallel sentences. We train a student model \hat{M} to mimic the sentence embeddings of the teacher M such that $\hat{M}(s_i) \approx M(t_i)$ and $\hat{M}(t_i) \approx M(s_i)$. The following mean squared loss function is minimized for each mini-batch β :

$$\frac{1}{|\beta|} \sum_{j \in \beta} \left[((M(s_j) - \hat{M}(s_j))^2 + (M(t_j) - \hat{M}(t_j))^2 \right]$$

Thus, the student \hat{M} learns to map each target and source sentence to the same location in vector space.

For the teacher M we compare two models:

1. all-mpnet-base-v2,² a model tuned for semantic search, trained on a large and diverse training set of 1B+ pairs (Denoted as mpnet).

2. sentence-t5-large,³ a T5 model tuned for sentence similarity tasks, trained on 2B pairs (Ni et al., 2022) (Denoted as st5).

Both above models have a final normalization layer which we remove prior to training to allow student model to learn the original vector space properties of the teacher model.




Figure 2: Translation search accuracy over training steps with _{grc}XLM-R student model.

We compare _{grc}mBERT and _{grc}XLM-R as the student model \hat{M} . We add a mean pooling layer and pair both student models with both teacher models (4 configurations) and train all the student parameters. With mpnet as the teacher, we train for 15 epochs, but with st5 the student model took twice as long to converge (See Figure 2), so we train for 30 epochs. We use a batch size of 128, a max sequence length of 128 tokens, 2000 warmup steps, and a learning rate of 2e-5. Every 500 training steps we measure STS performance as well as MSE loss and translation search accuracy on 5k hold-out pairs, keeping the model with best average performance across these tasks. Regardless of teacher model, _{grc}XLM-R took many more training steps to converge than _{grc}mBERT and was prone to catastrophic forgetting, which was alleviated by increasing the number of warmup steps.

We also experiment with training on parallel Modern Greek data from Wikipedia for 3 epochs and then on Ancient Greek data for 15 epochs if mpnet is the teacher and 6 and 30 epochs if st5 is the teacher. Although Modern Greek differs in many significant ways from Ancient Greek, training on this data gives the model additional exposure

²<https://huggingface.co/sentence-transformers/all-mpnet-base-v2>

³<https://huggingface.co/sentence-transformers/sentence-t5-large>

to aspects of Greek that have remained unchanged since antiquity, such as historical proper nouns. All evaluations are reported with and without training on this additional data.

2.3 Contrastive Learning

As a baseline against which to compare the models trained via the distillation method, we also train sentence embedding models using *Simple Contrastive Learning of Sentence Embeddings* (SimCSE), the contrastive learning method introduced by Gao et al. (2021). Contrastive learning pulls semantically-close neighbors together and pushes apart non-neighbors, and has been shown to be effective for training multilingual sentence embeddings (Gao et al., 2021; Tan et al., 2023). In addition to using dropout as noise, we use each Greek sentence and its English translation as positive pairs and other pairs in the same batch as negatives.

We use the CLS token representation and train for a maximum of 10 epochs with a batch size of 82, a max sequence length of 128 tokens, 2000 warmup steps, and a learning rate of 2e-5. Every 500 training steps we measure performance on the STS evaluation and translation search accuracy on the 5k hold-out pairs, keeping the highest performing model. As above, we also experiment with training on Modern Greek data for 3 epochs, and then Ancient Greek data for 10 epochs.

3 Training Data

3.1 Pre-training

Our pre-training dataset consists of the Ancient Greek text from the Perseus Digital Library⁴ and First1KGreek⁵, which are part of the Open Greek and Latin project.⁶ Different documents containing the same Greek work were removed. These sources contain approximately 32 million words of Ancient Greek text. Although Riemenschneider and Frank (2023) produced a much larger corpus of Greek text (100+ million words) using additional sources, at the time of writing their data is not publicly available. Our smaller dataset is sufficient for our purposes, as Reimers and Gurevych (2020) show that even languages with little pre-training in a multilingual student model can be effective targets for knowledge distillation.

⁴<https://github.com/PerseusDL/canonical-greekLit>

⁵<https://github.com/OpenGreekAndLatin/First1KGreek>

⁶<https://opengreekandlatin.org>

This dataset consist of Greek texts spanning a thousand years, covering different dialects and time periods of the language. We do not filter out any texts based on their dialect or time period.

In addition to the Greek text, we also collect all the English translations in the Open Greek and Latin project to finetune our models with an additional 10 million words of historical English text.

3.2 Preprocessing

Following Yamshchikov et al. (2022) and Singh et al. (2021), we lowercase all the Greek text and strip diacritics, but keep punctuation. Although diacritics contain important information for disambiguating between words that only differ by breathing marks or accent marks, the correct word can usually be inferred from context. The contextual nature of BERT models allows them to learn to use context to disambiguate.

3.3 Parallel Data

Human Aligned A portion of our parallel sentence dataset is taken from human aligned sources:

1. Verses of the Greek New Testament with English translations (15k pairs),
2. Verses of the Greek Septuagint with English translations (29k pairs),
3. Verses of the Greek works of Flavius Josephus with English translations (15k pairs),
4. Other minor sources: OPUS (Tiedemann and Nygaard, 2004), Greek Learner Texts⁷, manually aligned passages from Perseus and First1KGreek (total 23k pairs).

Translation Alignment The bulk of the parallel data is produced using translation alignment. We take all the texts from our pre-training corpus that have English translations and split them into sentences or sub-sentence segments (see Appendix B). We then use a two-step process to align Greek sentences with their English translations. First, we use Hunalign (Halász et al., 2007), a sentence-length heuristic and dictionary-based alignment technique on the translated texts. This produced an initial dataset of approximately 150k parallel sentences (including the human-aligned sources listed above).

Using this initial dataset, we trained a sentence embedding model with an aligned vector space for

⁷<https://greek-learner-texts.org>

English and Ancient Greek using SimCSE (See Section 2.3). Next, we use this model to align all the texts again, using a better alignment method introduced by Liu and Zhu (2023), dubbed Bertalign, which uses multilingual sentence embeddings to achieve state-of-the-art alignment quality. If the Greek and English documents are already aligned by sections, we align the sentences in each section individually. This increases alignment accuracy and makes it possible to keep the parts of the document that have good alignments and to discard the rest. Otherwise, if no section alignments exist, we run the aligner on the entire text.

We do not filter out multiple translations of the same Greek texts, since different translations can have different nuances and word choices, with the hope that the resulting sentence embeddings will be more robust to translation differences.

Finally, we remove all duplicate sentence pairs from the dataset and all pairs with very short sentences (<5 characters). We also ensure that no sentence pairs from the STS dataset (See Section 4.2) are included in the training data. This results in approximately 380k sentence pairs after holding out 5k pairs for evaluation purposes.

Modern Greek The Modern Greek (EL) sentence pairs from Wikipedia are taken from the OPUS project (Tiedemann and Nygaard, 2004). We remove all duplicate pairs and pairs with very short sentences (<10 characters), resulting in approximately 800k sentence pairs. This dataset contains a rich and diverse set of topics, including historical topics which will hopefully transfer to the Ancient Greek models. We compare all the models with and without training on this data.

4 Evaluations

4.1 Translation Similarity Search

The first measure of the quality of the sentence embeddings is each model’s accuracy at choosing the correct English translation for each Ancient Greek sentence from the 5k hold-out pairs. The score is computed as the percentage of sentence pairs for which the embedding of source sentence s_i has the closest cosine similarity to the embedding of translated sentence t_i out of all the target sentences. The accuracy is computed in both directions and averaged. The results are reported in Table 2.

The SimCSE models perform on this task better than the distillation models, which is not surpris-

Model	Accuracy
<i>SimCSE</i>	
GRCmBERT (GRC)	95.92
GRCmBERT (EL,GRC)	96.09
GRCXLM-R (GRC)	95.86
GRCXLM-R (EL,GRC)	96.64
<i>Teacher: sentence-t5-large</i>	
GRCmBERT (GRC)	87.78
GRCmBERT (EL,GRC)	90.80
GRCXLM-R (GRC)	87.02
GRCXLM-R (EL,GRC)	91.60
<i>Teacher: all-mpnet-base-v2</i>	
GRCmBERT (GRC)	87.77
GRCmBERT (EL,GRC)	89.15
GRCXLM-R (GRC)	86.48
GRCXLM-R (EL,GRC)	90.12

Table 2: Translation similarity search accuracy. Best result is bolded.

ing since they specifically trained to maximize the cosine similarity between translation pairs and minimize similarity between non-pairs. There is no significant difference in the performance between the two base models. All the models performed better when first trained on Modern Greek before Ancient Greek.

4.2 Semantic Textual Similarity

Sentence Pair	Score
Στωικοί ἀποφαίνονται σφαιροειδῆ τὸν κόσμον. Stoics declare the world to be spherical.	0.9
Στωικός νομίζει ὅτι ἡ γῆ σφαίρα ἐστιν. A Stoic thinks that the earth is a sphere.	
ἐπὶ δὲ τοῦ ὄρους τῇ ἀκρᾳ Διός ἐστιν ναός. On the top of the mountain is a temple of Zeus.	0.8
ὁ Ζεὺς οἰκεῖ ἐπὶ τὰ ὅρη ἐν Ὄλυμπῳ. Zeus dwells on the mountains in Olympus.	
Τὰ παιδία παίζουσιν ἐν τῇ ἀμμουδιᾷ. The children are playing in the sand.	0.5
Τὰ παιδία ἀναπαύονται ἐν τῷ κήπῳ. The children rest in the garden.	
Σωκράτης εἶδεν ἔξ βόας. Socrates saw six cows.	0.1
Ῥώμουλος εἶδεν ἔξ οἰωνοὺς ὅρνιθας. Romulus saw six birds of omen.	

Table 3: Example pairs from STS evaluation dataset. Scores are examples and not actual scores.

We compiled a dataset of Ancient Greek sentence pairs with gold scores to measure semantic textual similarity in the range [0,1], with 0 representing completely unrelated meaning, and 1 representing full semantic equivalence. Each sentence was given a corresponding English translation to

Model	GRC↔GRC	EN↔EN	GRC↔EN	Average
<i>SimCSE</i>				
GRCmBERT (GRC)	75.68	77.58	76.30	76.52
GRCmBERT (EL,GRC)	74.85	78.30	76.40	76.52
GRCXLM-R (GRC)	77.83	78.82	77.21	77.95
GRCXLM-R (EL,GRC)	78.27	79.11	77.76	78.38
<i>Teacher: sentence-t5-large</i>				
GRCmBERT (GRC)	82.17	87.54	84.02	84.58
GRCmBERT (EL,GRC)	84.84	89.33	86.37	86.84
GRCXLM-R (GRC)	82.37	85.37	82.56	83.43
GRCXLM-R (EL,GRC)	84.88	88.37	85.45	86.24
<i>Teacher: all-mpnet-base-v2</i>				
GRCmBERT (GRC)	82.30	87.60	84.68	84.86
GRCmBERT (EL,GRC)	84.84	88.77	86.28	86.63
GRCXLM-R (GRC)	83.80	87.07	84.53	85.13
GRCXLM-R (EL,GRC)	85.18	88.24	85.92	86.45

Table 4: Spearman rank correlation ρ between the cosine similarity of sentence embeddings and gold labels for STS dataset. Scores are reported as $\rho \times 100$. Best results are bolded. There are twice as many *GRC-EN* pairs as *GRC-GRC* pairs so their scores are not directly comparable.

allow for cross-lingual evaluation (See Table 3).

The gold scores for STS datasets are typically produced by averaging the scores from many human annotators. However, for Ancient Greek it is difficult to find enough annotators to produce high quality gold scores. Our solution is to use a Cross-Encoder⁸ to produce the gold scores based on the English translations of each pair. A Cross-Encoder takes two sentences as input and produces a similarity score in the range $[0, 1]$ without the need to encode the semantic properties of each sentence into a vector, and therefore performs better than cosine similarity between embeddings (See Figure 3). With this setup, we measure how closely each model can match the performance of the English Cross-Encoder. The accuracy of this method depends on how closely the English translations match the meaning of the Greek sentences. Therefore the English translations are reviewed by an expert to ensure that they are literal and accurate translations of the Greek text.

Due to the need to manually verify the translations for each pair, the STS dataset is relatively small. The dataset consists of 165 Ancient Greek sentences pairs, each having an English translation: $((a_{GRC}, a_{EN}), (b_{GRC}, b_{EN}))$. The GRC↔EN comparison can be performed two ways: $a_{GRC} \leftrightarrow b_{EN}$ and $a_{EN} \leftrightarrow b_{GRC}$ for a total of 330 GRC↔EN comparisons, 165 GRC↔GRC comparisons, and 165 EN↔EN comparisons.

⁸<https://huggingface.co/cross-encoder/stsb-roberta-base>




Figure 3: We use a Cross-Encoder (right) to produce STS gold scores which are used to evaluate our sentence embedding models, which are Bi-Encoders (left).

The score for each model is computed as Spearman correlation between gold scores and the cosine similarities between the sentence embeddings. The results are reported in Table 4.

The models trained via knowledge distillation significantly outperform the SimCSE models, showing that they have inherited the properties of the teacher models for STS tasks. The models with the *st5* teacher have a small lead, which is expected since *st5* was trained for STS tasks. All the models improve slightly when first trained on Modern Greek before Ancient Greek.

4.3 Semantic Retrieval

Sentence embeddings can be used for semantic retrieval tasks by ranking a set of passage embeddings by cosine similarity with a query embedding. Performing this process with our models on the Greek sentences in the Perseus and First1KGreek

corpora yields promising results. For example, the following query is correctly answered by several passages in the top 10 highest ranked passages:

Query: “Was Aristotle a student of Plato?”

- Ἀριστοτέλης Πλάτωνος μαθητής· οὗτος τὴν δι-
αλεκτικὴν συνεστήσατο. - Hyppolytus of Rome
Aristotle, a disciple of Plato — He established dialectics.
- ἀλλὰ καὶ τοῖς Πλάτωνος ἐγκαλέσαι ἂν τις δόγμασι
δι’ Ἀριστοτέλην, ἀποφοιτήσαντα τῆς διατριβῆς αὐ-
τοῦ ἐν καινοτομίᾳ. - Origen
But someone could also challenge certain doctrines of
Plato through Aristotle, who, upon completing his studies,
departed from his teachings with innovations.
- παρὸς Πλάτωνι Ἀριστοτέλης φιλοσοφήσας
μετελθὼν εἰς τὸ Λύκειον κτίζει τὴν Περιπατητικὴν
αἵρεσιν. - Clement of Alexandria
After studying philosophy under Plato, Aristotle, having
come to the Lyceum, founded the Peripatetic school.

To quantify the performance of each model for semantic retrieval, we compile a dataset of 40k Greek passages from the Perseus and First1KGreek corpora. We then produce 100 English queries (in the form of both phrases and questions) and associate them with relevant passages. We measure recall and mean average precision (mAP) for each model. The scores are reported in Table 5.

The SimCSE models perform poorly, which is expected since they were not trained for retrieval tasks. The models with the `mpnet` teacher, which was trained for semantic search, score highest by a large margin. The models with the `st5` teacher, which was trained for semantic textual similarity tasks, perform better than the SimCSE models but worse than the `mpnet` models. The models generally perform much better when trained on Modern Greek. Perhaps this is because many of the queries involve proper nouns for which Modern Greek data gave additional training examples, or perhaps the student models benefited from the additional English examples to learn the vector space properties of the teacher. The `GRCmBERT` models consistently perform better than `GRCXLM-R` models.

Overall performance on this task was rather poor even for the best models. An analysis of the top ranked passages for each query revealed that passages about related topics often ranked above the desired passages. In particular, it often confused proper names, e.g. preferring passages about other philosophers for queries about Plato.

4.4 Translation Bias

To determine whether the models are biased towards certain translation styles, especially those

Model	Recall@10	mAP@20
<i>SimCSE</i>		
GRCmBERT (GRC)	26.61	15.33
GRCmBERT (EL,GRC)	18.08	10.84
GRCXLM-R (GRC)	21.50	9.86
GRCXLM-R (EL,GRC)	29.56	15.08
<i>Teacher: sentence-t5-large</i>		
GRCmBERT (GRC)	41.34	25.37
GRCmBERT (EL,GRC)	49.63	36.17
GRCXLM-R (GRC)	34.88	20.07
GRCXLM-R (EL,GRC)	47.07	31.31
<i>Teacher: all-mpnet-base-v2</i>		
GRCmBERT (GRC)	63.60	44.97
GRCmBERT (EL,GRC)	69.87	53.00
GRCXLM-R (GRC)	53.84	36.42
GRCXLM-R (EL,GRC)	60.13	44.36

Table 5: Recall@10 and mAP@20 scores for English search queries and Ancient Greek passages. Best results are bolded.

included in the training set, a text with many different translations is needed. The New Testament is a good candidate for this, since many translations exist in different styles and eras of the English language. We take nine New Testament translations, ranging from literal (NASB), archaic (KJV), and paraphrase (MSG), all fully aligned at the verse level (7654 verses). There are no other Greek texts that we are aware of that have this many translations available for comparison. We generate embeddings for each verse from the Greek text and the translations. We also generate an “average” translation for each verse by averaging the embeddings of all the English translations. We take the cosine similarity between the Greek embedding and each translation and use it to compute the Mean Reciprocal Rank (MRR) across all verses, for each model:

$$MRR = \frac{1}{|T|} \sum_{v \in T} \frac{1}{rank_v}$$

where T is a set of verses in a translation and $rank_v$ is the rank of the translation for verse v . The results are reported in Table 6.

The literal translations score highest, and the score decreases the more non-literal the translations become, with the MSG translation having the lowest score. Surprisingly, the archaic KJV translation ranks highly, which is likely due to a high quantity of archaic English text in the training data. This suggests that the models are slightly biased to this older English translation style. Verses from two of the translations (NKJV and NET) were included in the training data. Despite being in the training data,

Model	KJV	NKJV*	NASB	ESV	RSV	NET*	NIV	NLT	MSG	Avg. Emb.
<i>SimCSE</i>										
GRCmBERT (GRC)	32.59	36.56	39.97	32.17	29.28	25.91	20.32	12.92	11.14	<u>52.04</u>
GRCmBERT (EL,GRC)	33.27	36.05	40.79	31.97	28.74	26.17	20.15	12.79	11.21	<u>51.75</u>
GRCXLM-R (GRC)	35.78	37.85	38.17	32.15	29.76	26.03	20.58	13.14	11.32	<u>48.11</u>
GRCXLM-R (EL,GRC)	35.34	36.74	38.02	32.96	30.06	25.76	20.75	13.09	11.25	<u>48.93</u>
<i>Teacher: sentence-t5-large</i>										
GRCmBERT (GRC)	29.63	30.70	30.13	27.81	25.90	23.05	20.09	14.43	12.49	<u>78.66</u>
GRCmBERT (EL,GRC)	28.73	30.66	29.98	28.02	25.82	23.39	19.70	13.93	12.24	<u>80.42</u>
GRCXLM-R (GRC)	31.82	29.70	29.14	28.10	26.39	23.82	20.38	14.24	12.90	<u>76.40</u>
GRCXLM-R (EL,GRC)	30.41	30.01	29.08	28.35	26.82	23.61	19.75	13.68	12.36	<u>78.82</u>
<i>Teacher: all-mpnet-base-v2</i>										
GRCmBERT (GRC)	31.76	31.15	29.93	30.04	28.94	23.33	19.52	13.93	11.98	<u>72.32</u>
GRCmBERT (EL,GRC)	30.42	31.20	30.37	30.35	28.68	23.51	19.53	13.76	11.81	<u>73.26</u>
GRCXLM-R (GRC)	37.25	31.31	29.91	30.00	29.42	23.32	19.67	13.97	12.32	<u>65.74</u>
GRCXLM-R (EL,GRC)	33.08	31.21	29.64	30.25	29.59	23.55	19.32	13.61	11.89	<u>70.76</u>

* Verses from the NET and NKJV were included in parallel training data.

Table 6: Mean Reciprocal Rank (MRR) $\times 100$ of cosine similarity between Greek verses of the New Testament and English translations, as well as MRR of per-verse averaged embedding of all the translations. Highest translation MRR for each model is bolded. MRR of averaged embedding is underlined if it is higher than any of the translations.

there does not appear to be bias to the NET since it consistently ranks lower than other translations. The NKJV ranks highly, but does not consistently outrank other literal translations. Interestingly, the average embedding of all the translations ranked highest by a significant margin.

5 Discussion and Future Work

Overall, the base models mBERT and XLM-R performed similarly except for the semantic retrieval task where the mBERT-derived models have a sizeable lead. The reason for this is unclear, since these models have different tokenizers, parameter counts, and vocabularies. It is also unclear how much the pre-training process affects the results. **An area of future research would be to investigate the effect of student model architecture, tokenizer, and pre-training on the ability of the student model to learn from the teacher model.**

The main limitation of using multilingual knowledge distillation to train sentence embedding models is that the embeddings produced are almost entirely derived from English translations, which could be undesirable if the goal is to study Ancient Greek text without any prior translator’s interpretation. Furthermore, the student model can never fully replicate the performance of the teacher model when transferring to another language, since translated sentences are often not entirely semantically equivalent to their source sentences, especially when removed from the original context via

translation alignment.

Although contamination from modern languages is not a big concern for the tasks in this paper, **there could be issues of anachronisms when searching Ancient Greek texts with English**. Furthermore, using texts from such a **long chronological period of the Greek language could introduce additional lexical polysemy as Greek words changed in meaning over time**. This could explain why the averaged embedding of many translations had a higher MRR than any individual translation source in Table 6, since the combination of many translations represents a higher degree of polysemy. **In future work, such historical polysemy could be measured by sampling translations of words from texts of different historical periods**. This could help to determine whether the high MRR of the averaged embedding is a useful result or simply an artifact of a potentially high amount of polysemy in the training data.

6 Conclusion

In this paper, we have shown that multilingual knowledge distillation is an effective approach for training sentence embedding models for Ancient Greek, in spite of the lack of available training data compared to modern, high-resource languages. In addition, we have produced a new dataset of parallel Ancient Greek and English sentences as well as evaluation datasets for translation search, semantic textual similarity, and semantic retrieval, which we make publicly available.

References

- Steven Bird, Ewan Klein, and Edward Loper. 2009. *Natural Language Processing with Python*, 1st ed edition. O’Reilly, Beijing ; Cambridge [Mass.].
- Alexis Conneau, Kartikay Khandelwal, Naman Goyal, Vishrav Chaudhary, Guillaume Wenzek, Francisco Guzmán, Edouard Grave, Myle Ott, Luke Zettlemoyer, and Veselin Stoyanov. 2020. **Unsupervised cross-lingual representation learning at scale**. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 8440–8451, Online. Association for Computational Linguistics.
- Charlie Cowen-Breen, Creston Brooks, Johannes Haubold, and Barbara Graziosi. 2023. **Logion: Machine Learning for Greek Philology**.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. **BERT: Pre-training of deep bidirectional transformers for language understanding**. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota. Association for Computational Linguistics.
- Tianyu Gao, Xingcheng Yao, and Danqi Chen. 2021. **SimCSE: Simple contrastive learning of sentence embeddings**. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 6894–6910, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Péter Halácsy, Andras Kornai, Viktor Nagy, László Németh, and Viktor Trón. 2007. **Parallel corpora for medium density languages**. In *Recent Advances in Natural Language Processing IV*, pages 247–258.
- Valentin Hofmann, Janet Pierrehumbert, and Hinrich Schütze. 2021. **Superbizarre is not superb: Derivational morphology improves BERT’s interpretation of complex words**. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 3594–3608, Online. Association for Computational Linguistics.
- Kyle P. Johnson, Patrick J. Burns, John Stewart, Todd Cook, Clément Besnier, and William J. B. Mattingly. 2021. **The Classical Language Toolkit: An NLP framework for pre-modern languages**. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing: System Demonstrations*, pages 20–29, Online. Association for Computational Linguistics.
- John Koutsikakis, Ilias Chalkidis, Prodromos Malakasiotis, and Ion Androutsopoulos. 2020. **GREEK-BERT: The Greeks visiting Sesame Street**. <https://arxiv.org/abs/2008.12014v2>.
- Lei Liu and Min Zhu. 2023. **Bertalign: Improved word embedding-based sentence alignment for Chinese–English parallel corpora of literary texts**. *Digital Scholarship in the Humanities*, 38(2):621–634.
- Jianmo Ni, Gustavo Hernandez Abrego, Noah Constant, Ji Ma, Keith Hall, Daniel Cer, and Yinfei Yang. 2022. **Sentence-t5: Scalable sentence encoders from pre-trained text-to-text models**. In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 1864–1874, Dublin, Ireland. Association for Computational Linguistics.
- Hyunji Hayley Park, Katherine J. Zhang, Coleman Hally, Kenneth Steimel, Han Liu, and Lane Schwartz. 2021. **Morphology matters: A multilingual language modeling analysis**. *Transactions of the Association for Computational Linguistics*, 9:261–276.
- Telmo Pires, Eva Schlinger, and Dan Garrette. 2019. **How multilingual is multilingual BERT?** In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 4996–5001, Florence, Italy. Association for Computational Linguistics.
- Nils Reimers and Iryna Gurevych. 2019. **Sentence-BERT: Sentence embeddings using Siamese BERT-networks**. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3982–3992, Hong Kong, China. Association for Computational Linguistics.
- Nils Reimers and Iryna Gurevych. 2020. **Making monolingual sentence embeddings multilingual using knowledge distillation**. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 4512–4525, Online. Association for Computational Linguistics.
- Frederick Riemenschneider and Anette Frank. 2023. **Exploring Large Language Models for Classical Philology**.
- Pranaydeep Singh, Gorik Rutten, and Els Lefever. 2021. **A pilot study for BERT language modelling and morphological analysis for ancient and medieval Greek**. In *Proceedings of the 5th Joint SIGHUM Workshop on Computational Linguistics for Cultural Heritage, Social Sciences, Humanities and Literature*, pages 128–137, Punta Cana, Dominican Republic (online). Association for Computational Linguistics.
- Weiting Tan, Kevin Heffernan, Holger Schwenk, and Philipp Koehn. 2023. **Multilingual representation distillation with contrastive learning**. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics*, pages 1477–1490, Dubrovnik, Croatia. Association for Computational Linguistics.

Jörg Tiedemann and Lars Nygaard. 2004. The OPUS corpus - parallel and free: <http://logos.uio.no/opus>. In *Proceedings of the Fourth International Conference on Language Resources and Evaluation (LREC'04)*, Lisbon, Portugal. European Language Resources Association (ELRA).

Ivan Yamshchikov, Alexey Tikhonov, Yorgos Pantis, Charlotte Schubert, and Jürgen Jost. 2022. BERT in plutarch’s shadows. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*, pages 6071–6080, Abu Dhabi, United Arab Emirates. Association for Computational Linguistics.

A Appendix: Training Details

Parameter	GRCmBERT	GRCXLM-R
Batch Size	140	128
Learning Rate	2e-5	2e-5
LR Scheduler	linear	linear
Epochs	10	10
Warmup Steps	2000	2000
Mask Percentage	15%	15%

Table 7: Pre-training hyperparameters

Parameter	GRCmBERT	GRCXLM-R
Batch Size	128	128
Learning Rate	2e-5	2e-5
LR Scheduler	linear	linear
Max Seq. Length	128	128
Pooling	mean	mean
Embedding Dim.	768	768
<i>Teacher: all-mpnet-base-v2</i>		
Epochs (GRC)	15	15
Epochs (EL)	3	3
GRC Warmup Steps	2000	2000
EL Warmup Steps	2000	8000
<i>Teacher: sentence-t5-large</i>		
Epochs (GRC)	30	30
Epochs (EL)	6	6
GRC Warmup Steps	2000	2000
EL Warmup Steps	2000	2000

Table 8: Knowledge distillation hyperparameters

Parameter	GRCmBERT	GRCXLM-R
Batch Size	82	82
Learning Rate	2e-5	2e-5
LR Scheduler	linear	linear
Warmup Steps	2000	2000
Max Seq. Length	128	128
Epochs (GRC)	10	10
Epochs (EL)	3	3
Pooling	CLS	CLS
Embedding Dim.	768	768

Table 9: SimCSE hyperparameters

B Appendix: Sentence Segmentation

For translation alignment, it is not necessary that each segment be a sentence, since the alignment process can handle 1-many, many-1 or many-to-many relations. The Greek texts in our corpus contain punctuation, so we segment them by period (.), question mark (;), and raised dot (·). Some of the Greek texts use a colon instead of a raised dot, and in these cases we treat colons as raised dots. For the English texts we first segment using the NLTK sentence tokenizer (Bird et al., 2009) then further subdivide these segments by semicolon (;) and colon (:).